



# **Comparative Analysis of Temporal Convolutional Networks for Object Classification**

Valli Chidambaram, Justin Davis, Paris Dinh, Amari Hoogland,  
Sofie Lange, Kamen Shah, Lei Teng, & Zhengwu Yuan

# Overview

## Problem

Traditional object classification models exhibit intrinsic problems when inferring on video sequences

## Goal

Design a model that can account for previous classifications in the hope to reduce temporal instability in classifying objects within videos



# Applications



Surveillance Systems



Face Detection



Self-driving Vehicles



Sports Analysis



Autonomous Systems




Traffic analysis / optimization

# Approach

## 1. Dataset

Obtain a dataset with annotated bounding boxes for video sequences



## 2. Implement Base Network

Run YOLO V1 to establish baseline

## 3. Integrate Time-Dependent layers

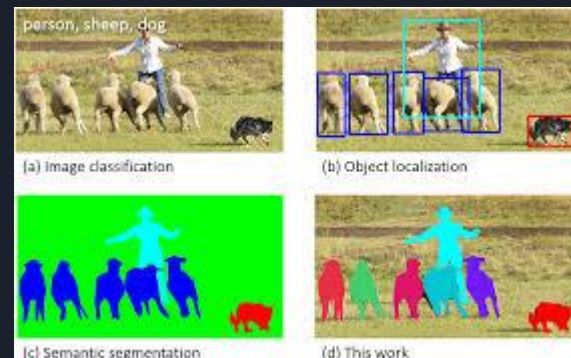

- RNN
- LSTM
- TCNN

## 4. Hyperparameter Tuning



# Dataset Selection

- Originally looking at vehicle datasets only
- Many datasets contain classified vehicles in static, stand-alone images, e.g. KITTI, COCO, Stanford Cars
- Difficult to find high-quality, publicly available labelled training data if we limit the scope to only vehicles
- Settled on YouTube Bounding Boxes



# Dataset

YouTube Bounding Boxes - augmented with classical object tracking from OpenCV



Class labels include - bicycle, bus, dog, cat, truck, car, train, potted plant, person, etc.

- Bounding boxes and class labels for ~240,000 publicly available YouTube videos
- 15-20s video segments with b.b. at 1 fps
- 23 different classes
- features objects in natural settings without editing or post-processing
- recording quality similar to a hand-held cell phone camera

# Problems with Dataset

- Unbalanced classes, far more people than anything else
- Bounding boxes from tracking are inaccurate, sometimes disappear
- Most videos only have one labeled object, but several unlabeled objects
  - Loss increases when predicting unlabeled boxes during training
  - Precision decreases when predicting unlabeled boxes






# Models

YOLO v1 - Baseline	RNN	LSTM	TCNN	Attention
<ul style="list-style-type: none"><li>• Convolutional neural network</li><li>• Predicts bounding boxes and class probabilities</li><li>• YOLO is extremely fast (45 frames per second)</li></ul>	<ul style="list-style-type: none"><li>• Recurrent Neural network</li><li>• Exhibits temporal dynamic behavior.</li></ul>	<ul style="list-style-type: none"><li>• Long Short Term Memory</li><li>• Deals with the vanishing gradient problem</li></ul>	<ul style="list-style-type: none"><li>• Temporal Convolutional Neural Network</li><li>• Convolutional architecture that combines simplicity, autoregressive prediction, and very long memory for video sequences</li></ul>	<ul style="list-style-type: none"><li>• Mechanism based off the concept of directing focus to certain factors with data</li><li>• Usually used bidirectionally on text data</li></ul>

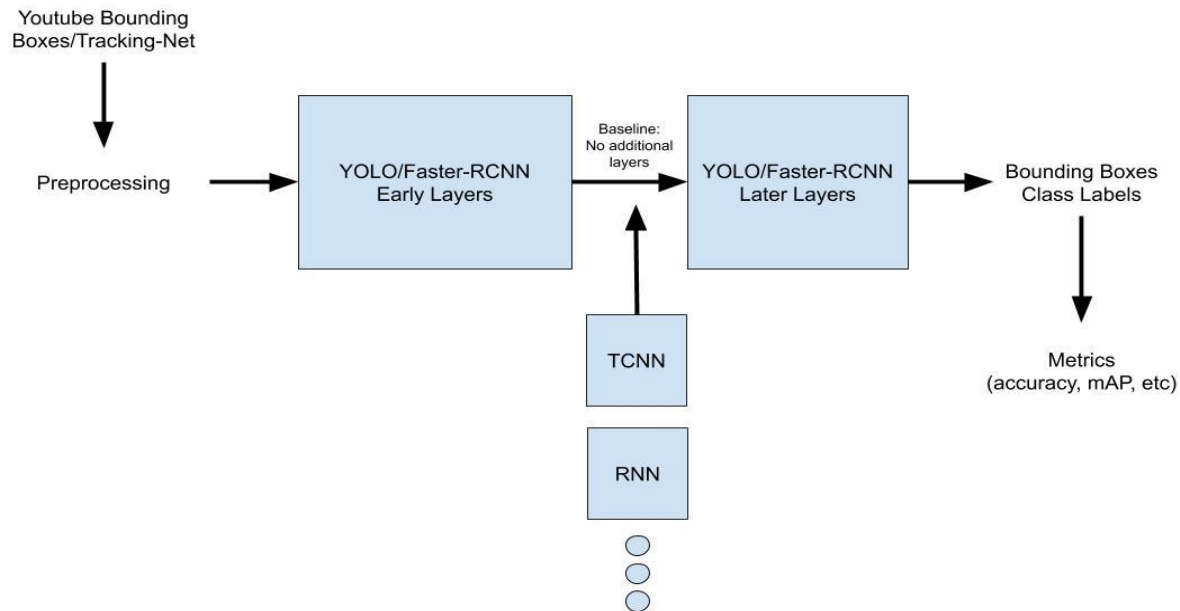


# YOLO V1



It consists of 24 convolutional layers where the first 20 convolutional layers followed by a maxpool layer and the last 4 convolutional layers followed by 2 fully connected layer.

# Architecture Diagram

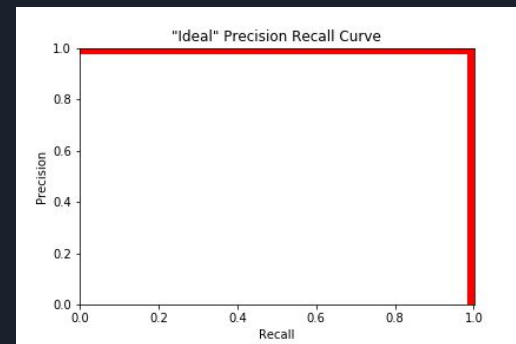

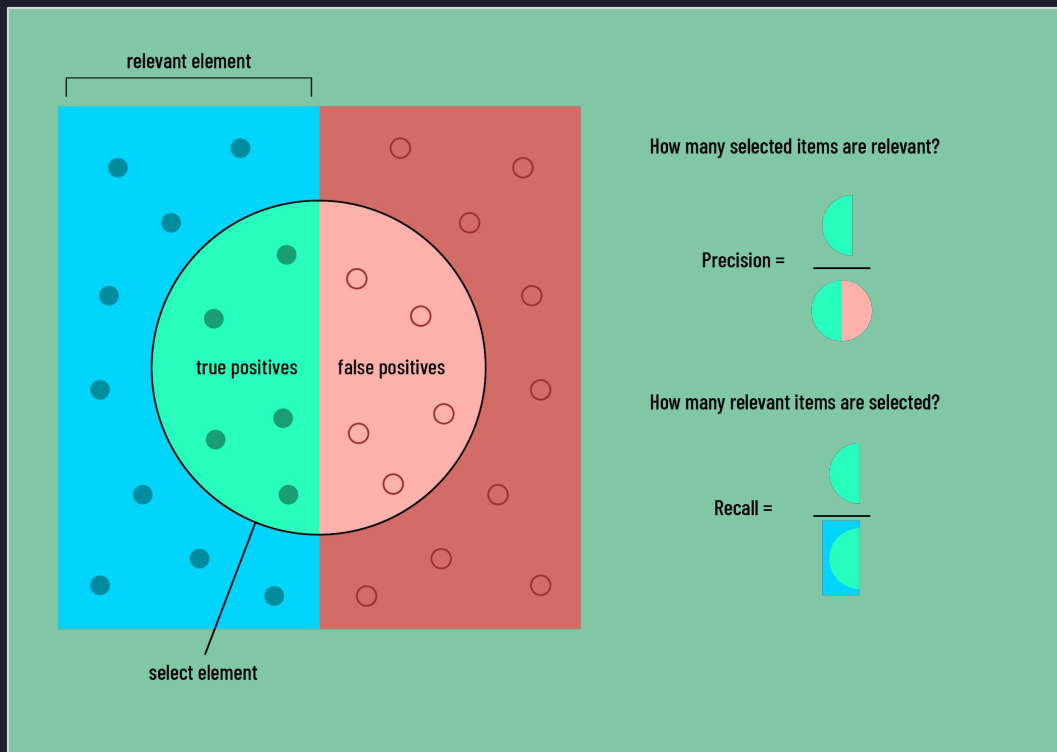


# Results

- Individual Model Metrics
  - Recall v. Precision
  - Mean Average Precision Score (mAP)
- Comparative Results
- Visualizations

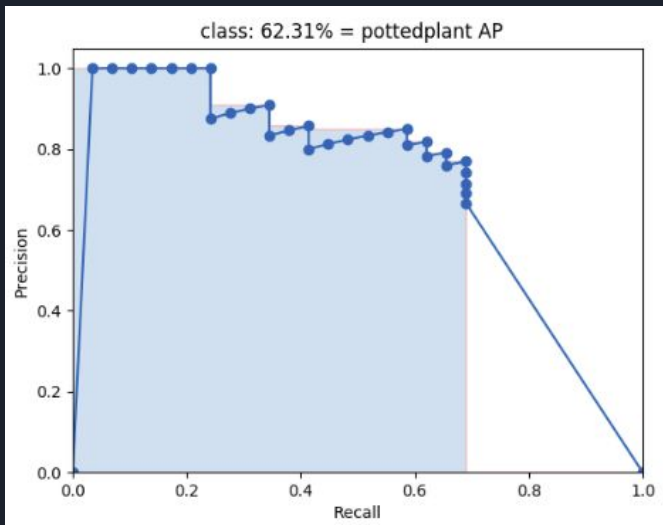


# Precision, Recall, and IOU



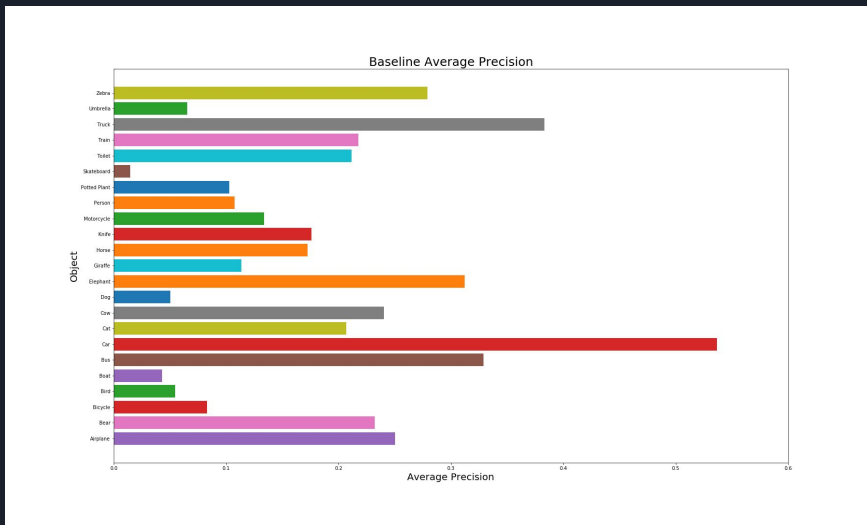
# mAP (Mean Avg. Precision) Score

- Commonly used metric for object detection algorithm comparison
- Used to evaluate performance in popular object detection challenges such as [MS COCO](#) and [PASCAL VOC](#)



Example graph from <https://github.com/Cartucho/mAP>

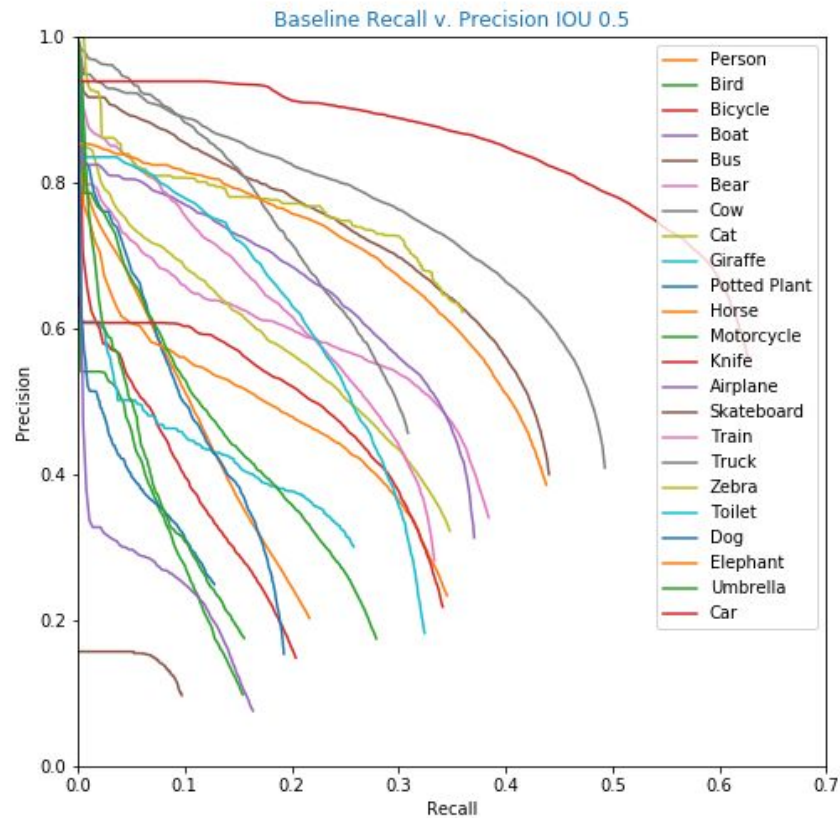
AP is the area under this curve **per class**



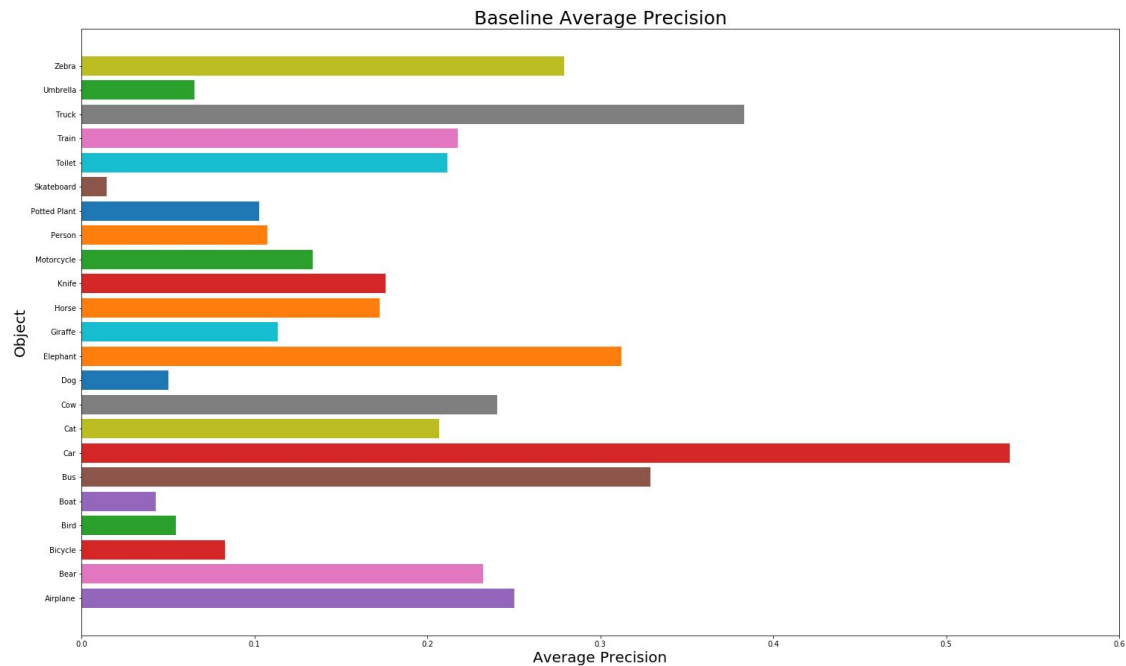
mAP is the average of AP scores for all classes

# YOLO v1

- Baseline

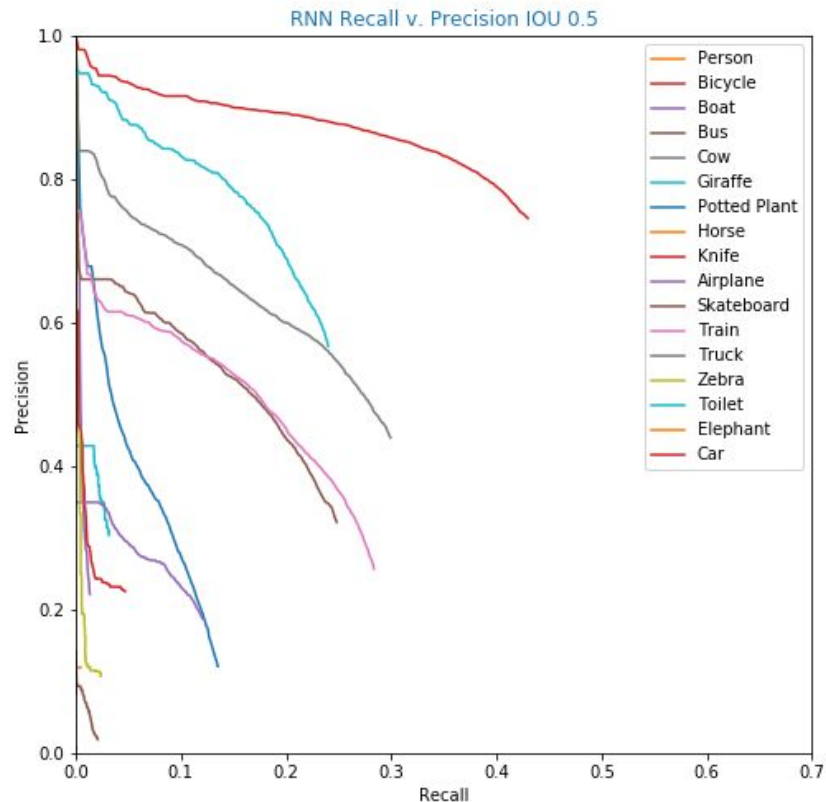


# YOLO v1



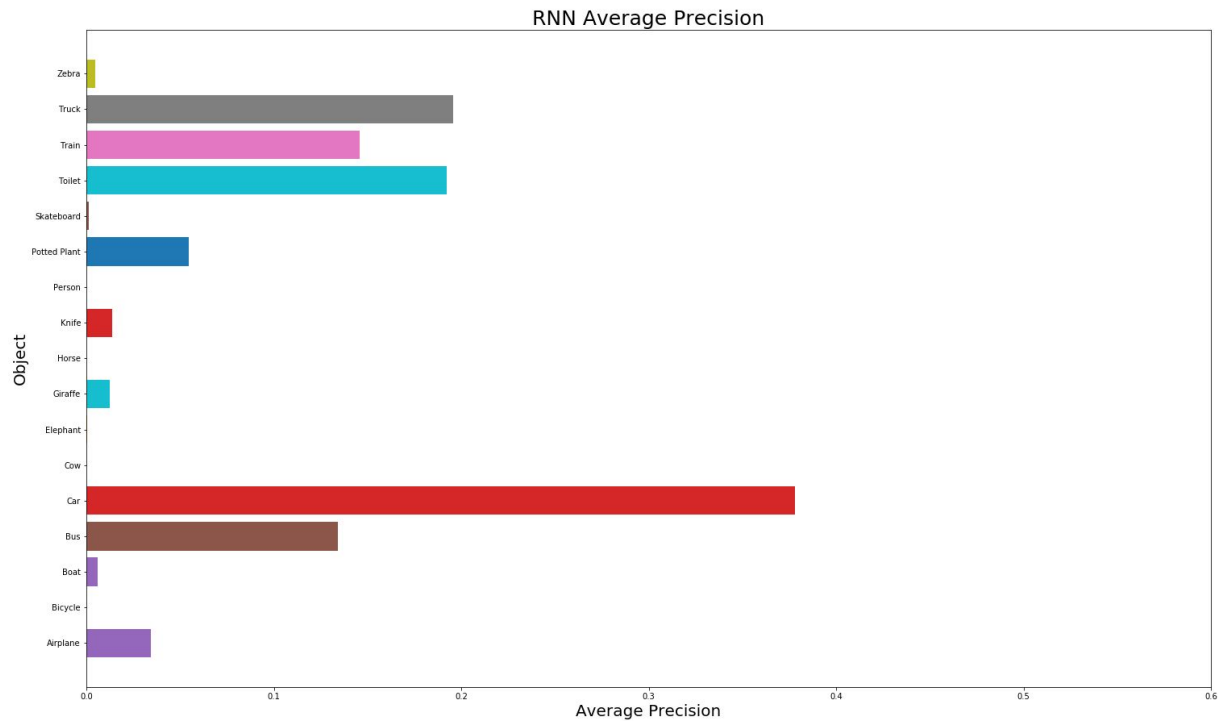
# RNN

- Hard to train, loss wouldn't decrease
- Lots of tuning to get it to work at all



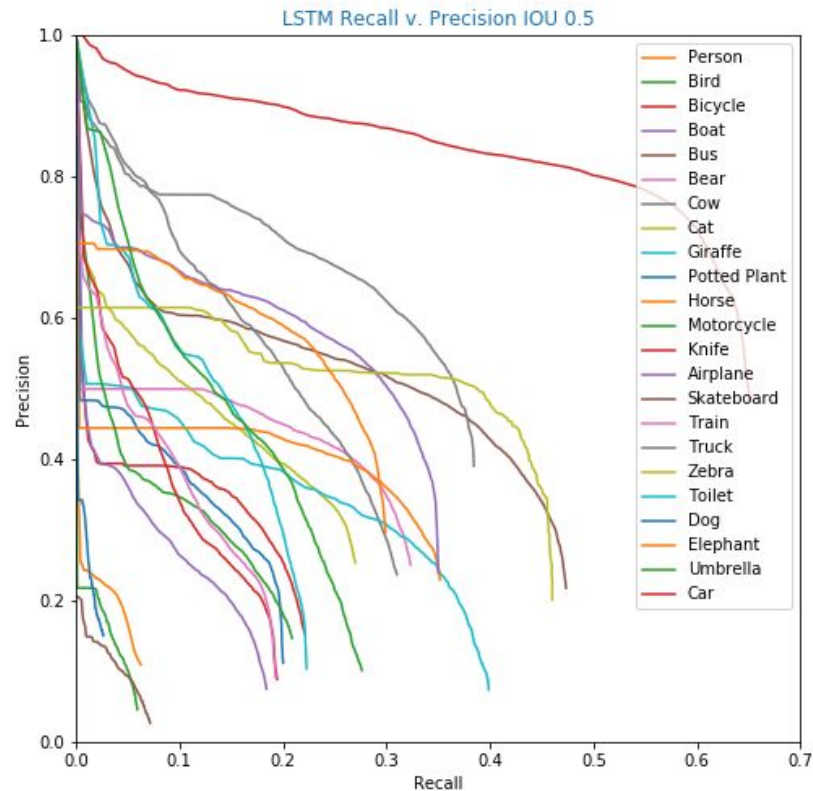


# RNN

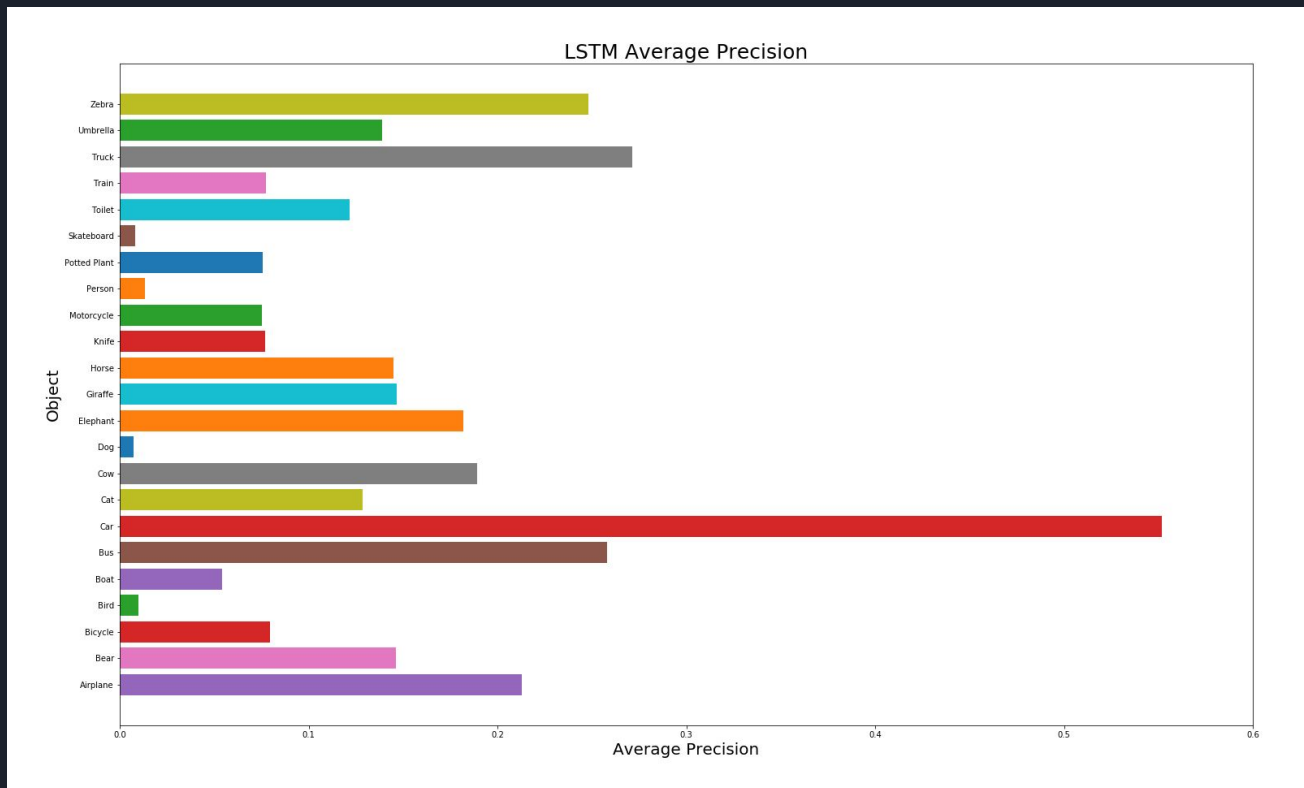


# LSTM

- Best performance of time-dependent models

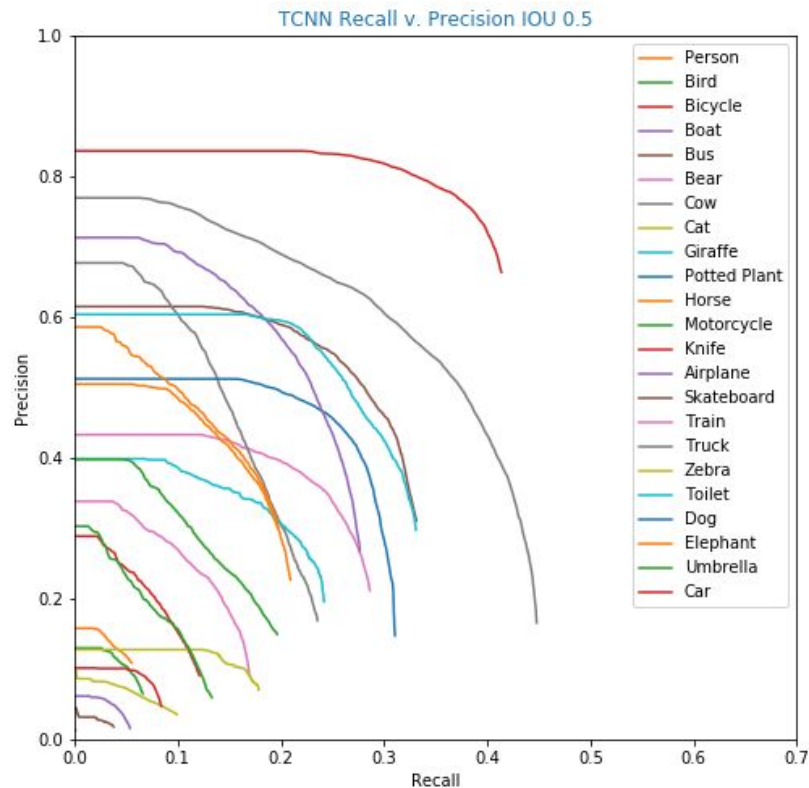


# LSTM

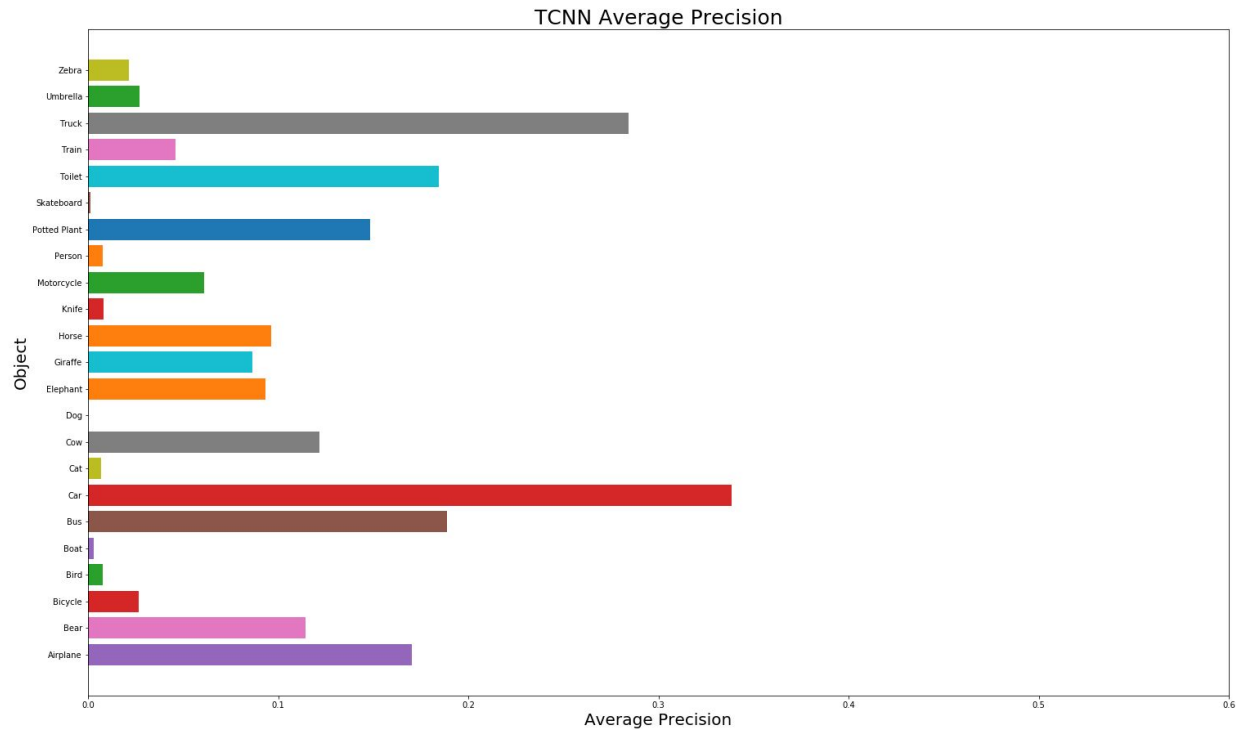


# TCNN

- Only trained once
- No changes to architecture
- No tuning
- Surprisingly well performing

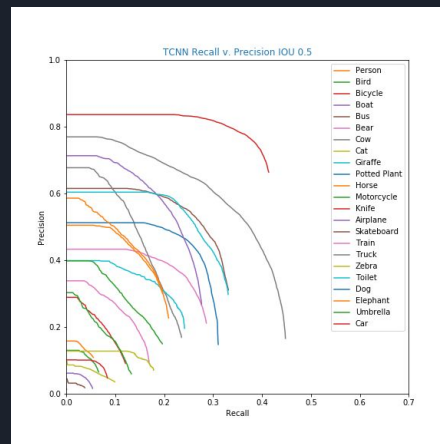
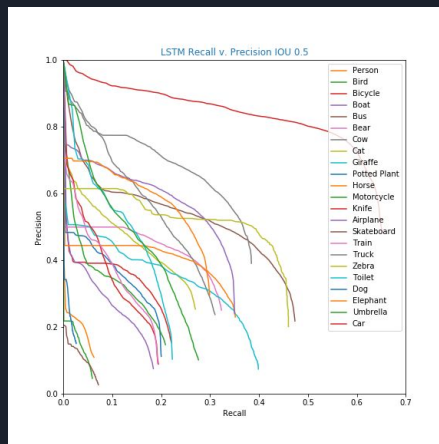

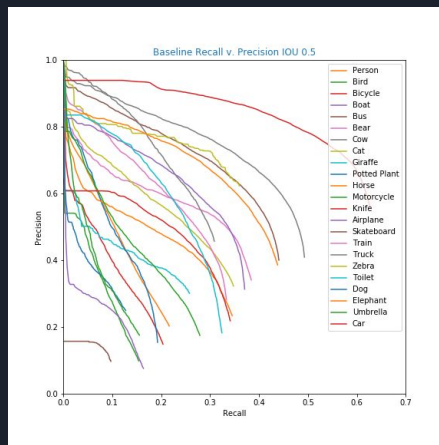


# TCNN



# Comparative Results

- All recurrent models did worse than the baseline
- Recurrent models were harder to train
- LSTM was the best recurrent model
- TCNN had the least hyperparameter tuning and testing



# Results

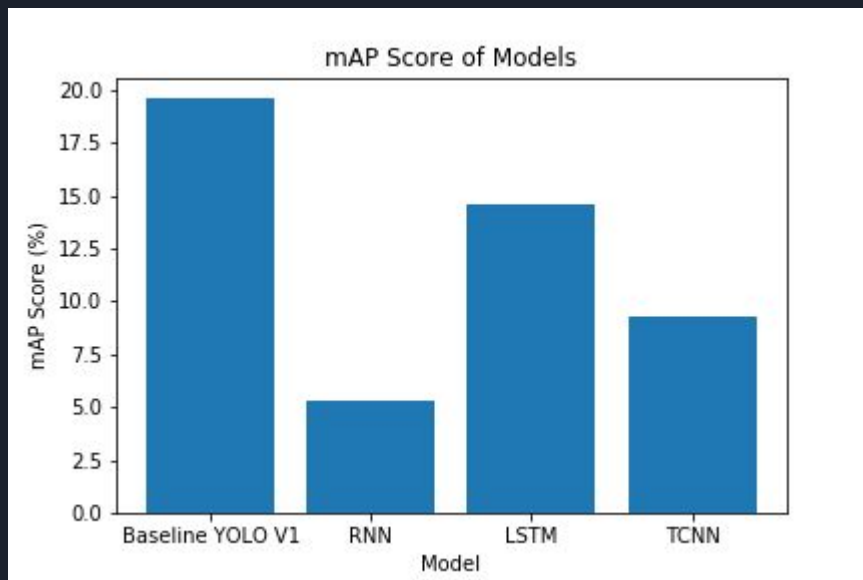
IOU 0.5 mAP Scores:

Baseline YOLO: 19.62%


RNN: 5.34%

LSTM: 14.63%


TCNN: 9.29%



# YOLO V1 Baseline








LSTM



TCNN



RNN



LSTM



TCNN



RNN


# YOLO V1 Baseline








LSTM



TCNN



RNN



# Conclusion

- **Recurrent models can significantly improve temporal stability for classification**
  - They do not directly improve overall mAP scores
  - LSTM models achieve the best balance between mAP scores and temporal stability
  - TCNN works well with little to no tuning
- **Models can be further improved**
  - Better convolutional network
  - Further hyperparameter tuning
  - Different combinations of recurrent layers

# Challenges

- Temporal models require immense amount of memory leading to out-of-memory errors on GPU
- Unbalanced dataset, almost half the videos are people
- Inaccurate dataset, bounding boxes propagated by classical tracking algorithms aren't always accurate
  - Sometimes bounding boxes disappear during video
  - Almost every video only has one object labeled
- Training time
- Difficulty debugging neural networks
- Overfitting
- Hyperparameter tuning is difficult





# Future Work

- Run RNN and LSTM over all features in a sequence
- Use random sequence length inputs while training recurrent models
- Grid Search over all parameters in validation to find optimal inferencing parameters
- Implement recurrent layers within a larger convolutional neural network like Yolo v3
- Research further the effects of messy data (heterogeneous values, missing entries, and large errors)



Questions?





# Thank You

David Motta

Christian Butterfield

Northrop Grumman